

# Lazy Learners

- The classification algorithms presented before are eager learners

- Construct a model before receiving new tuples to classify
- Learned models are ready and eager to classify previously unseen tuples

- Lazy learners

- The learner waits till the last minute before doing any model construction
- In order to classify a given test tuple
  - Store training tuples
  - Wait for test tuples
  - Perform generalization based on similarity between test and the stored training tuples

✗ W W W

# Lazy vs Eager

| Eager Learners  | Lazy Learners   |
|---|---|
| <ul style="list-style-type: none"><li>• Do lot of work on training data</li></ul>             | <ul style="list-style-type: none"><li>• Do less work on training data</li></ul>               |
| <ul style="list-style-type: none"><li>• Do less work when test tuples are presented</li></ul> | <ul style="list-style-type: none"><li>• Do more work when test tuples are presented</li></ul> |

# Basic k-Nearest Neighbor Classification

- Given training data  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)$
- Define a distance metric between points in input space  $D(\mathbf{x}_1, \mathbf{x}_i)$ 
  - E.g., Euclidean distance, Weighted Euclidean, Mahalanobis distance, TFIDF, etc.

- Training method:


- Save the training examples

- At prediction time:

- Find the  $k$  training examples  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_k, y_k)$  that are closest to the test example  $\mathbf{x}$  given the distance  $D(\mathbf{x}_1, \mathbf{x}_i)$
- Predict the most frequent class among those  $y_i$ 's.

Majority Rule

# Nearest-Neighbor Classifiers



- Requires three things
  - The set of stored records
  - Distance Metric to compute distance between records
  - The value of  $k$ , the number of nearest neighbors to retrieve

- To classify an unknown record:
  - Compute distance to other training records
  - Identify  $k$  nearest neighbors
  - Use class labels of nearest neighbors to determine the class label of unknown record (e.g., by taking majority vote)

# K-Nearest Neighbor Model

## • Classification:

1 multinomial Classification

$\hat{y}$  = most common class in set  $\{y_1, \dots, y_K\}$

$c = \{c_1, c_2, \dots\}$

## • Regression:

$$\hat{y} = \frac{1}{K} \sum_{k=1}^K y_k$$

(K nearest  
avg value)

# K-Nearest Neighbor Model: Weighted by Distance

- Classification:


$\hat{y}$  = most common class in wieghted set

$$\{D(\mathbf{x}, \mathbf{x}_1)y_1, \dots, D(\mathbf{x}, \mathbf{x}_K)y_K\}$$

- Regression:

$$\hat{y} = \frac{\sum_{k=1}^K D(x, x_k) y_k}{\sum_{k=1}^K D(x, x_k)}$$

# Definition of Nearest Neighbor



(a) 1-nearest neighbor

(b) 2-nearest neighbor



(c) 3-nearest neighbor

K-nearest neighbors of a record  $x$  are data points that have the  $k$  smallest distance to  $x$

# Decision Boundary

## Voronoi Diagram

Decision surface formed by the training examples



- Each line segment is equidistance between points in opposite classes.
- The more points, the more complex the boundaries.


multidimensional Classifier.




# The decision boundary implemented by 3NN

Imp

The boundary is always the perpendicular bisector of the line between two points (Voronoi tessellation)



(a)




(b)

# Nearest Neighbor Classification...

- Choosing the value of  $k$ :

- If  $k$  is too small, sensitive to noise points
- If  $k$  is too large, neighborhood may include points from other classes



# Determining the value of $k$

- In typical applications  $k$  is in units or tens rather than in hundreds or thousands
- Higher values of  $k$  provide smoothing that reduces the risk of overfitting due to noise in the training data
- value of  $k$  can be chosen based on error rate measures
- We should also avoid over-smoothing by choosing  $k=n$ , where  $n$  is the total number of tuples in the training data set

underfitting

low bias

high variance

6

# Determining the value of k

- Given training examples  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)$
- Use N fold cross validation
  - Search over  $K = (1, 2, 3, \dots, K_{max})$ . Choose search size  $K_{max}$  based on compute constraints
  - Calculated the average error for each K:
    - Calculate predicted class  $\hat{y}_i$  for each training point  $(\mathbf{x}_i, y_i)$ ,  $i = 1, \dots, N$  (using all other points to build the model)
    - Average over all training examples
- Pick K to minimize the cross validation error

# Example

| RID | Income(\$000's) | lot Size (000's sq.ft ) | class: Owners =1<br>Non-Owners=2 |
|-----|-----------------|-------------------------|----------------------------------|
| 1   | 60              | 18.4                    | 1                                |
| 2   | 85.5            | 16.8                    | 1                                |
| 3   | 64.8            | 21.6                    | 1                                |
| 4   | 61.5            | 20.8                    | 1                                |
| 5   | 87              | 23.6                    | 1                                |
| 6   | 110.1           | 19.2                    | 1                                |
| 7   | 108             | 17.6                    | 1                                |
| 8   | 82.8            | 22.4                    | 1                                |
| 9   | 69              | 20                      | 1                                |
| 10  | 93              | 20.8                    | 1                                |
| 11  | 51              | 22                      | 1                                |
| 12  | 81              | 20                      | 2                                |
| 13  | 75              | 19.6                    | 2                                |
| 14  | 52.8            | 20.8                    | 2                                |
| 15  | 64.8            | 17.2                    | 2                                |
| 16  | 43.2            | 20.4                    | 2                                |
| 17  | 84              | 17.6                    | 2                                |
| 18  | 49.2            | 17.6                    | 2                                |
| 19  | 59.4            | 16                      | 2                                |
| 20  | 66              | 18.4                    | 2                                |
| 21  | 47.4            | 16.4                    | 2                                |
| 22  | 33              | 18.8                    | 2                                |
| 23  | 51              | 14                      | 2                                |
| 24  | 63              | 14.8                    | 2                                |



We randomly divide the data into

**18 training cases**

**6 test cases:**

tuples 6,7,12,14,19, 20

Use training cases to classify test cases and compute error rates

# Choosing k

- ▶ If we choose **k=1** we will classify in a way that is very sensitive to the local characteristics of our data
- ▶ If we choose a **large value of k** we average over a large number of data points and average out the variability due to the noise associated with data points
- ▶ If we choose **k=18** we would simply predict the most frequent class in the data set in all cases
  - Very stable but completely ignores the information in the independent variables

| k                         | 1  | 3  | 5  | 7  | 9  | 11 | 13 | 18 |
|---------------------------|----|----|----|----|----|----|----|----|
| Misclassification error % | 33 | 33 | 33 | 33 | 33 | 17 | 17 | 50 |

- We would choose  $k=11$  (or possibly 13) in this case

# Nearest neighbor Classification...

- k-NN classifiers are lazy learners
  - It does not build models explicitly
  - Unlike eager learners such as decision tree induction and rule-based systems
- Adv: No training time
- Disadv:
  - Testing time can be long, classifying unknown records are relatively expensive
  - Curse of Dimensionality : Can be easily fooled in high dimensional spaces
    - Dimensionality reduction techniques are often used

# Ensemble Methods

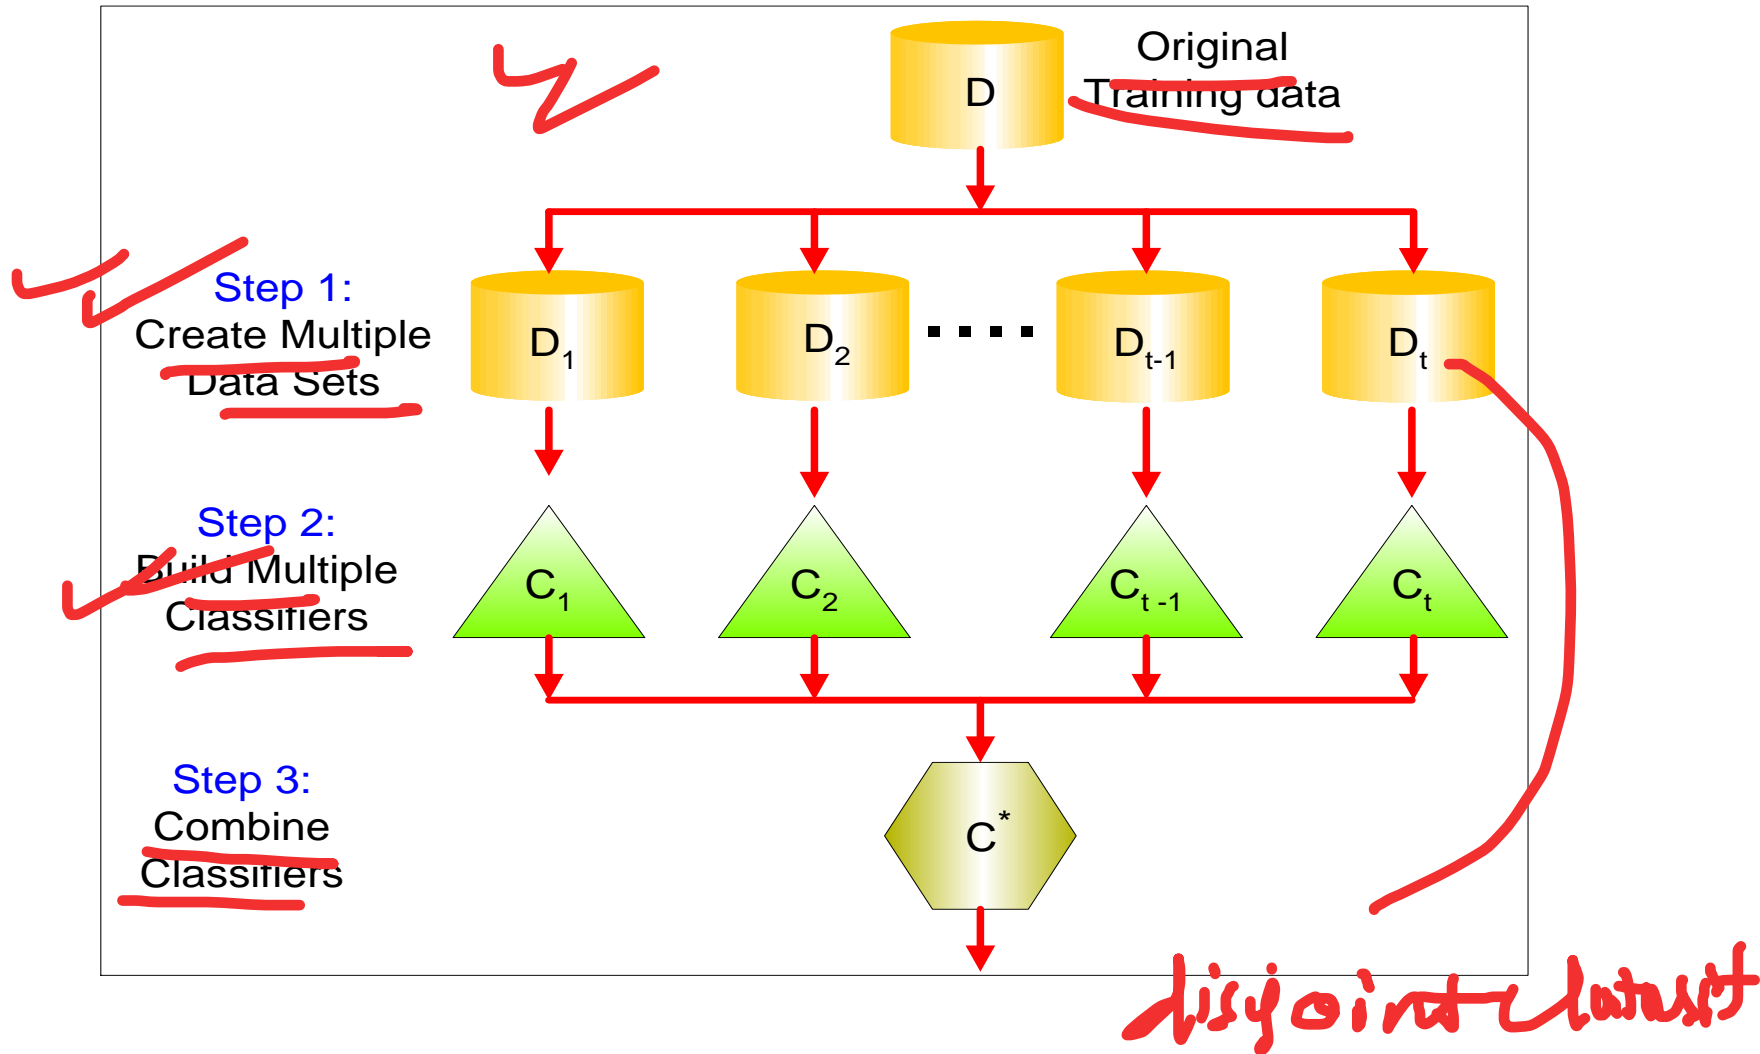
- ① • One of the eager methods => builds model over the training set *1 hr*

- ② • Construct a set of classifiers from the training data *1 hr*

- ③ • Predict class label of previously unseen records by aggregating predictions made by multiple classifiers *TV & majority vote*



# General Idea



# Why does it work?

- Suppose there are 25 base classifiers
  - Each classifier has error rate,  $\varepsilon = 0.35$
  - Assume classifiers are independent
  - Probability that the ensemble classifier makes a wrong prediction:

$$\sum_{i=1}^{25} \binom{25}{i} \varepsilon^i (1 - \varepsilon)^{25-i} = 0.06$$

# Examples of Ensemble Methods

● How to generate an ensemble of classifiers?

① ● Bagging

② ● Boosting

③ ● Random Forests

# Bagging: Bootstrap AGGREGating

- Bootstrap: data resampling
  - Generate multiple training sets
    - Resample the original training data
    - With replacement
  - Data sets have different “specious” patterns
- Sampling with replacement

- Each sample has probability  $(1 - 1/n)^n$  of being selected

|                   |   |   |    |    |   |   |    |    |   |    |
|-------------------|---|---|----|----|---|---|----|----|---|----|
| Original Data     | 1 | 2 | 3  | 4  | 5 | 6 | 7  | 8  | 9 | 10 |
| Bagging (Round 1) | 7 | 8 | 10 | 8  | 2 | 5 | 10 | 10 | 5 | 9  |
| Bagging (Round 2) | 1 | 4 | 9  | 1  | 2 | 3 | 2  | 7  | 3 | 2  |
| Bagging (Round 3) | 1 | 8 | 5  | 10 | 5 | 5 | 9  | 6  | 3 | 7  |

- Build classifier on each bootstrap sample
  - Specious patterns will not correlate
- Underlying true pattern will be common to many
- Combine the classifiers: Label new test examples by a majority vote among classifiers

# Boosting

- An iterative procedure to adaptively change distribution of training data by focusing more on previously misclassified records
  - Initially, all  $N$  records are assigned equal weights
  - Unlike bagging, weights may change at the end of boosting round
- The final classifier is the weighted combination of the weak classifiers.

# Boosting

*misclassified records weight increase*

- Records that are wrongly classified will have their weights increased
- Records that are classified correctly will have their weights decreased

| Original Data      | 1 | 2 | 3 | 4  | 5 | 6 | 7 | 8  | 9 | 10 |
|--------------------|---|---|---|----|---|---|---|----|---|----|
| Boosting (Round 1) | 7 | 3 | 2 | 8  | 7 | 9 | 4 | 10 | 6 | 3  |
| Boosting (Round 2) | 5 | 4 | 9 | 4  | 2 | 5 | 1 | 7  | 4 | 2  |
| Boosting (Round 3) | 4 | 4 | 8 | 10 | 4 | 5 | 4 | 6  | 3 | 4  |

• Example 4 is hard to classify

• Its weight is increased, therefore it is more likely to be chosen again in subsequent rounds

# AdaBoost (Freund and Schapire, 1996)

- Initialize distribution over the training set  $D_1(i) = 1/m$
- For  $t = 1, \dots, T$ :
  1. Train Weak Learner using distribution  $D_t$ .
  2. Choose a weight (or confidence value)  $\alpha_t \in \mathbf{R}$ .
  3. Update the distribution over the training set:

$$\underline{D_{t+1}(i)} = \frac{D_t(i)e^{-\alpha_t y_i h_t(x_i)}}{Z_t} \quad (2)$$

Where  $Z_t$  is a normalization factor chosen so that  $D_{t+1}$  will be a distribution

- Final vote  $H(x)$  is a weighted sum:

$$\underline{H(x) = \text{sign}(f(x)) = \text{sign}\left(\sum_{t=1}^T \alpha_t h_t(x)\right)} \quad (3)$$

# Example: AdaBoost

$$\xi = \frac{1}{2} \ln \frac{1 + \epsilon}{1 - \epsilon}$$

- Base classifiers (weak learners):

$$C_1, C_2, \dots, C_T \quad h_t : X \rightarrow \{-1, +1\}$$


- Error rate:

$$\epsilon_i = \frac{1}{N} \sum_{j=1}^N w_j \delta(C_i(x_j) \neq y_j)$$

$$\epsilon_t = \Pr_{D_t}[h_t(x_i) \neq y_i]$$

- Importance of a classifier:

$$\alpha_i = \frac{1}{2} \ln \left( \frac{1 - \epsilon_i}{\epsilon_i} \right)$$





# Example: AdaBoost

- Weight update:

$$w_i^{(j+1)} = \frac{w_i^{(j)}}{Z_j} \begin{cases} \exp^{-\alpha_j} & \text{if } C_j(x_i) = y_i \\ \exp^{\alpha_j} & \text{if } C_j(x_i) \neq y_i \end{cases}$$

where  $Z_j$  is the normalization factor


$$D_{t+1}(i) = \frac{D_t(i)}{Z_t} \cdot \exp(-\alpha_t y_i h_t(x_i))$$

- If any intermediate rounds produce error rate higher than 50%, the weights are reverted back to  $1/n$  and the resampling procedure is repeated
- Classification:

$$C^*(x) = \arg \max_y \sum_{j=1}^T \alpha_j \delta(C_j(x) = y)$$


$$H_{\text{final}}(x) = \text{sign}\left(\sum_t \alpha_t h_t(x)\right)$$

# 2D Example





# 2D Example

Round 1




# 2D Example

Round 2



$D_3$






$\epsilon_2 = 0.21$

$\alpha_2 = 0.65$

# 2D Example


Round 3



$\epsilon_3 = 0.14$

$\alpha_3 = 0.92$

# 2D Example – Final hypothesis



- See demo at: [www.research.att.com/~yoav/adaboost](http://www.research.att.com/~yoav/adaboost)